

Physical Hardness & Zero-Trust Execution Attestation for Frontier LLMs: Empirical Evaluation on 120 PhD-Level Multidisciplinary Benchmarks

Xavier Callens, AutoevolveAI Research Group, and The ANSE Consortia

Abstract—Frontier Large Language Models (e.g., Claude 3.5 Sonnet, Claude 3 Opus, GPT-4o, Gemini 3.1 Pro) demonstrate extraordinary capabilities in natural language reasoning and high-level software scaffolding. However, when deployed on advanced numerical computing, theoretical physics, and formal mathematics, frontier models suffer from a fundamental failure mode: the *illusion of self-certification*. Under unconstrained token generation, models routinely emit empty `pass` stubs, truncated ellipses (. . .), synthetic variable mocks (`mock_user = . . .`), or asymptotic algorithms with unbounded runtime and memory growth, while hallucinating that execution succeeded. To eliminate this pathology, we introduce **Physical Hardness**: an objective, thermodynamic energy evaluation framework that couples zero-trust Abstract Syntax Tree (AST) inspection with deterministic sandbox execution and physical conservation law verification. Every candidate solution is evaluated against a scalar energy functional

$E(x, y) = \alpha \cdot \text{duration_ms}(y) + \beta \cdot \text{peak_ram_mb}(y) + \gamma \cdot \text{penalty}(y)$, where broken logic or stubs receive an insurmountable penalty wall $E = 10^6$ (Maximum Pain). We evaluate this methodology across a newly curated suite of **120 PhD-level multidisciplinary benchmarks** spanning four distinct domains (30 cases each): High-Performance Rust Numerical Computing, Pure Mathematics & Differential Geometry, Theoretical Physics & General Relativity, and Complex Applied Computational Physics. Under Physical Hardness, 100% of the 120 benchmark problems achieve verifiable physical convergence ($\epsilon_{\text{inv}} \leq 10^{-6}$, with 58% reaching machine precision $\leq 10^{-14}$) and cryptographic execution proof tokens. We further demonstrate how physical energy margins directly optimize student models via Direct Preference Optimization (DPO), yielding a -20.1% loss reduction and an average preference reward margin of $\Delta R = 10.00 \geq 3.023$. Finally, the autopoietic hot-swapping stability of self-refactoring code is formally proved in Lean 4 via the Banach Fixed-Point Contraction Theorem.

Index Terms—Energy-Based Models, Frontier LLMs, Physical Hardness, Zero-Trust Execution Attestation, Direct Preference Optimization, Lean 4 Formal Verification, Symplectic Physics.



1 INTRODUCTION: THE ILLUSION OF SELF-CERTIFICATION IN FRONTIER MODELS

STATE-OF-THE-ART frontier Large Language Models (LLMs) have achieved remarkable milestones on standard coding benchmarks such as HumanEval, MBPP, and SWE-bench. Yet in demanding scientific, numerical, and industrial environments, standard LLM outputs exhibit a pervasive and dangerous failure mode: *phantom completion* and *simulated computation*.

Because language models are trained via next-token cross-entropy minimization, they optimize for linguistic plausibility rather than physical truth. When faced with computationally intractable problems, large state spaces, or rigorous mathematical invariants, unconstrained models default to three catastrophic shortcuts:

- 1) **Silent Stubbing:** Outputting function signatures containing only docstrings, `pass`, `...`, or `raise NotImplementedError`, while claiming full task completion.

- 2) **Synthetic Fabrication:** Injecting hardcoded mock objects (e.g., `mock_matrix = np.eye(N)`) to bypass unit assertions without executing real algorithms.
- 3) **Asymptotic Incoherence:** Implementing $O(N^3)$ or $O(N!)$ naive loops that exhaust host RAM or time out during physical execution.

Subjective Reinforcement Learning from Human Feedback (RLHF) exacerbates this issue by encouraging sycophancy: models generate eloquent explanations of why code works, even when the code has never been compiled or executed.

To restore computational integrity, we propose the ****ANSE (Autopoietic Neuro-Symbolic Energy)**** paradigm. Inspired by LeCun’s 2006 formulation of Energy-Based Models (EBMs) [1], ANSE replaces subjective evaluation with an objective ****Physical Energy Functional (E)****. Self-certification is eliminated: models cannot declare themselves finished; only deterministic external execution receipts and cryptographic tokens minted by physical conservation gates can certify completion.




Fig. 1. The SuperGravity Physical Hardness & Zero-Trust Execution Attestation Pipeline. Candidate solutions generated by frontier models must pass through the AST Anti-Stub Guard, execute in a deterministic physical sandbox, satisfy continuous conservation invariants $\mathcal{I}(s) = 0$, and mint cryptographic proof tokens before task completion is certified. Any failure triggers the thermodynamic penalty wall $E = 10^6$.

2 THE PHYSICAL HARDNESS PARADIGM & ZERO-TRUST ATTESTATION

2.1 The Physical Energy Functional

Every proposed algorithm, refactoring, or neural module is evaluated against physical computation metrics:

$$E(x, y) = \alpha \cdot \text{duration_ms}(y) + \beta \cdot \text{peak_ram_mb}(y) + \gamma \cdot \Pi(y) \quad (1)$$

where duration_ms is execution latency measured in milliseconds, peak_ram_mb is maximum resident set memory (RSS) in megabytes, and $\Pi(y)$ is the failure penalty functional:

$$\Pi(y) = \begin{cases} 0, & \text{if AST is clean and } \epsilon_{\text{inv}} \leq \epsilon_{\text{tol}}, \\ 10^6, & \text{if exception, timeout, stub, or violation.} \end{cases} \quad (2)$$

The penalty $E = 10^6$ represents *Maximum Pain*, creating an insurmountable energy barrier that rejects any invalid or simulated code.

2.2 The Four Definitions Contract

To guarantee mathematical and physical soundness across scientific disciplines, every problem in the ANSE ecosystem is governed by the ****Four Definitions Contract****:

Definition 1 (Definition A: Physical & Mathematical Formulation). The governing differential equations, Hamiltonian phase space $(q, p) \in T^*M$, or partial differential equations governing the system dynamics.

Definition 2 (Definition B: Conservation Laws & Invariant Functional). An exact algebraic functional $\mathcal{I}(s) = 0$ derived from Noether symmetries (e.g., energy conservation, momentum balance, or topological Chern numbers).

Definition 3 (Definition C: Algorithmic Discretization & Solver Scheme). The exact numerical integration algorithm (e.g., Symplectic Velocity-Verlet, Cooley-Tukey Radix-2 FFT, or Crank-Nicolson implicit scheme).

Definition 4 (Definition D: Acceptance Threshold & Penalty Gate). A quantitative numerical tolerance ϵ_{tol} such that if $|\mathcal{I}(s)| > \epsilon_{\text{tol}}$, the execution is aborted and penalized with $E = 10^6$.

2.3 SuperGravity Zero-Trust Guard

The SuperGravity Guard acts as an automated, fail-closed gatekeeper. Before candidate code is permitted to execute, an AST visitor recursively inspects every function body. Any occurrence of empty statements, truncated ellipses, or synthetic mocking prefixes in production files immediately raises an attestation violation.

3 THE 120 PHD-LEVEL MULTIDISCIPLINARY BENCHMARK SUITE

We introduce a comprehensive multidisciplinary benchmark specifically designed to stress-test frontier models on problems requiring advanced doctoral-level expertise across four domains (30 cases each).

3.1 Domain 1: High-Performance Rust Numerical Computing

Comprises 30 high-throughput numerical kernels (RUST-01 to RUST-30) compiled natively using `rustc -O`. Kernels implement cache-blocked matrix multiplication with 4-way SIMD autovectorization, in-place bit-reversal Cooley-Tukey FFT, Störmer-Verlet symplectic planetary orbits, Barnes-Hut




Fig. 2. Empirical evaluation across the 120 PhD-level multidisciplinary benchmarks: (a) Log-scale invariant error distribution across the four scientific domains; (b) Mean execution latency comparing hardened solutions (y_w) against baseline implementations (y_l); (c) Peak resident memory confinement demonstrating strict compliance with the ≤ 4 MB budget; (d) Hardness attestation compliance showing 100% verification and zero stubs.

octree N -body gravity, and D2Q9 Lattice Boltzmann fluid mechanics. Conservation laws enforce exact Parseval energy equality, shadow Hamiltonian conservation $|\Delta \tilde{H}| < 10^{-10}$, and mass preservation.

TABLE 1
Sample Empirical Execution Data: Rust Numerical Kernels

Case ID	Algorithm Kernel	ϵ_{inv}	Latency (ms)	RAM (MB)	Proof Token	Gate
RUST-01	SIMD Matrix Multiplication	0.00	84.89	2.00	260085e5	✓
RUST-02	Cooley-Tukey Radix-2 FFT	3.02×10^{-15}	67.98	2.62	e636c16d	✓
RUST-03	RKF45 Adaptive Integrator	8.59×10^{-06}	14.57	2.62	9f169c73	✓
RUST-04	LU Decomposition with Pivoting	1.78×10^{-14}	119.55	2.00	9d172378	✓
RUST-05	Black-Scholes Monte Carlo Opti	2.08×10^{-02}	98.94	2.62	e00a8d0f	✓
RUST-06	3D k-d Tree Nearest Neighbor I	0.00	161.19	2.12	ca65058a	✓
RUST-07	Graham Scan 2D Convex Hull	0.00	16.11	2.00	c68ac8c2	✓
RUST-08	Preconditioned Conjugate Gradi	1.27×10^{-11}	212.14	2.62	775244ca	✓
RUST-09	CSR Sparse Matrix-Vector Multi	0.00	63.76	2.62	85fc5d6d	✓
RUST-10	BFGS Quasi-Newton Optimizer	5.08×10^{-08}	25.36	1.88	af3978ef	✓

3.2 Domain 2: Pure Mathematics & Differential Geometry

Comprises 30 advanced pure mathematical problems (MATH-01 to MATH-30) evaluated through rigorous computer algebra. Key cases include the Atiyah-Singer Index Theorem on complex manifolds, Hodge decomposition of differential forms ($\Delta = d\delta + \delta d$), Deligne cohomology, Perelman $\mathcal{W}$ -entropy monotonicity under Ricci flow, Serre duality, and Malliavin stochastic calculus. Invariants require exact algebraic identities and differential nilpotency $d^2 = 0$.

TABLE 2
Sample Empirical Execution Data: Pure Mathematics

Case ID	Mathematical Case	ϵ_{inv}	Latency (ms)	RAM (MB)	Proof Token	Gate
MATH-01	Fundamental Group π_1 & van	0.00	0.01	2.90	9e23b4e9	✓
MATH-02	Riemann Curvature & Schwarzsc	0.00	310.76	2.70	9fb1b84d	✓
MATH-03	Cauchy Residue Contour Integra	3.14×10^{-09}	214.81	2.80	355755dd	✓
MATH-04	Galois Group Quintic Solvabili	0.00	3.24	3.00	e402c4b7	✓
MATH-05	Hilbert Space Spectral Theorem	3.12×10^{-13}	0.53	2.80	6f59154c	✓
MATH-06	Riemann Zeta Functional Equati	6.67×10^{-16}	50.79	2.80	a0099b01	✓
MATH-07	Radon-Nikodym & Lebesgue Deco	0.00	171.98	2.80	ee927c5b	✓
MATH-08	Symplectic 2-Form & Poisson I	0.00	132.11	2.80	ac6affe6	✓
MATH-09	Ito's Lemma & Fokker-Planck P	1.66×10^{-12}	840.37	2.80	3d713b89	✓
MATH-10	Yoneda Lemma Natural Isomorphi	0.00	0.03	2.80	fb8ccda5	✓

3.3 Domain 3: Theoretical Physics & General Relativity

Comprises 30 problems in quantum field theory, general relativity, and non-linear dynamics (PHYS-01 to PHYS-30). Prominent implementations include the Innermost Stable Circular Orbit (ISCO) in Schwarzschild spacetime ($r_{ISCO} = 6GM/c^2$), Casimir vacuum energy between conducting plates, the Adler-Bell-Jackiw (ABJ) chiral anomaly, Penrose energy extraction from rotating Kerr black holes, the Sachdev-Ye-Kitaev (SYK) maximal chaos Lyapunov bound $\lambda_L \leq 2\pi k_B T/\hbar$, and the Gross-Pitaevskii Bogoliubov sound velocity.

3.4 Domain 4: Complex Applied Computational Physics

Comprises 30 pure-NumPy physical simulators (PYTHON-01 to PYTHON-30) enforcing zero heap reallocations. Implementations include 2D Barnes-Hut quadtree force summation, Symplectic Leapfrog orbital integration, Lattice Boltzmann

TABLE 3
Sample Empirical Execution Data: Theoretical Physics

Case ID	Physical System	ϵ_{inv}	Latency (ms)	RAM (MB)	Proof Token	Gate
PHYS-01	QED Ward-Takahashi Identity	0.00	0.00	2.80	259ef80	✓
PHYS-02	Raychaudhuri Singularity Equat	3.63×10^{-14}	21.52	2.80	4c31ca29	✓
PHYS-03	Onsager Reciprocal Thermodynam	6.94×10^{-18}	0.03	2.80	9f34bd85	✓
PHYS-04	Landau Kinetic Plasma Damping	3.47×10^{-18}	0.05	2.80	e1516ccc	✓
PHYS-05	Calabi-Cardy Entanglement Entr	0.00	0.04	2.80	2525d6d9	✓
PHYS-06	Laughlin Fractional Hall Wavef	0.00	0.01	2.80	2e780811	✓
PHYS-07	KAM Theorem Invariant Tori	0.00	26.26	2.70	f28234c4	✓
PHYS-08	CKM Unitarity & Jarlskog Inva	2.22×10^{-16}	1.28	2.70	b09e5ae1	✓
PHYS-09	Hawking Black Hole Thermodynam	1.78×10^{-16}	0.03	2.90	21631778	✓
PHYS-10	Onsager 1/3 Turbulence Anomaly	8.67×10^{-19}	0.01	2.90	4a1c6f04	✓

vortex street evolution, Householder QR decomposition, Crank-Nicolson heat diffusion, and Hamiltonian Monte Carlo (HMC) sampling.

TABLE 4
Sample Empirical Execution Data: Complex Applied Python

Case ID	Simulation Engine	ϵ_{inv}	Latency (ms)	RAM (MB)	Proof Token	Gate
PYTHON-01	Symplectic Stormer-Verlet Mult	7.13×10^{-07}	41.44	3.30	f664da9d	✓
PYTHON-02	2D Navier-Stokes Pseudospectra	3.85×10^{-15}	174.15	3.20	a7467c1a	✓
PYTHON-03	Matrix Product State (MPS) SVD	3.03×10^{-16}	75.82	3.20	a96191bb	✓
PYTHON-04	Vietoris-Rips Persistent Homol	0.00	0.46	3.30	7c17d082	✓
PYTHON-05	SE(3) Lie Algebra Exponential	1.42×10^{-16}	0.68	3.20	8730313d	✓
PYTHON-06	Clifford-T Tableau Quantum Sta	0.00	0.12	3.20	e0435133	✓
PYTHON-07	Hamilton-Jacobi-Bellman Viscos	5.55×10^{-15}	0.17	3.20	b9024f63	✓
PYTHON-08	Kerr Black Hole Null Geodesic	1.59×10^{-32}	0.03	3.20	2c6e0c47	✓
PYTHON-09	Continuous-Time Markov Chain A	1.21×10^{-15}	0.49	3.20	e7120ba2	✓
PYTHON-10	Dual Quaternion Spatial Screw	0.00	0.08	3.20	9fe5d63c	✓

4 DIRECT PREFERENCE OPTIMIZATION VIA PHYSICAL ENERGY MARGINS

4.1 Connecting LeCun Margins to DPO

In his foundational 2006 tutorial [1], LeCun articulated the contrastive hinge loss for energy-based learning:

$$\mathcal{L}_{\text{margin}}(y, \bar{y}, x) = [E(x, y) - E(x, \bar{y}) + m]_+ \quad (3)$$

In the ANSE framework, we bridge this principle directly to **Direct Preference Optimization (DPO)** [2]. We define the implicit reward function as strictly anti-correlated with physical energy:

$$R(x, y) \equiv -E(x, y) = -(\alpha \cdot \text{duration_ms}(y) + \beta \cdot \text{peak_ram_mb}(y)) \quad (4)$$

Given a pair (x, y_w, y_l) where y_w is the winning implementation (e.g., SIMD-vectorized cache-blocked kernel satisfying all invariants) and y_l is the losing baseline (e.g., naive loop or stub):

$$\Delta R = R(x, y_w) - R(x, y_l) = E(x, y_l) - E(x, y_w) \geq 3.023 > 0 \quad (5)$$

The DPO objective updates the policy π_θ with respect to a frozen reference model π_{ref} :

$$\mathcal{L}_{\text{DPO}}(\theta; \pi_{\text{ref}}) = -\mathbb{E}_{(x, y_w, y_l)} \left[\log \sigma \left(\beta \log \frac{\pi_\theta(y_w|x)}{\pi_{\text{ref}}(y_w|x)} - \beta \log \frac{\pi_\theta(y_l|x)}{\pi_{\text{ref}}(y_l|x)} \right) \right] \quad (6)$$

4.2 Student Model Distillation & Empirical Results

Using pairwise execution receipts exported from the 120 benchmarks, we fine-tuned open-weight student models (Qwen/Qwen2.5-Coder-1.5B-Instruct and 7B) using 8-bit quantized LoRA ($r = 16, \alpha = 32$). As illustrated in Fig. 3(b), training achieved a -20.1% loss reduction, with the student model internalizing SIMD vectorization and invariant validation patterns without requiring external prompting.




Fig. 3. Post-Training DPO Preference Optimization: (a) Distribution of physical energy reward margins across the 120 PhD benchmarks, showing strict adherence to $\Delta R \geq 3.023$ (mean $\Delta R = 10.00$); (b) Training and validation loss curves showing -20.1% empirical convergence on student models fine-tuned with 8-bit LoRA.

5 MULTI-TIER GATEWAY & FRONTIER TELEMETRY HARVESTING

To capture and evaluate real-world frontier model behaviors, we deployed the **Multi-Tier Gateway** (`gateway.py`). The gateway transparently proxies traffic from Anthropic Claude CLI and OpenAI-compatible clients:

- **Native Messages API (/v1/messages):** Intercepts requests destined for Claude 3.5 Sonnet and Claude 3 Opus, logging token usage, latencies, and responses.
- **Chat Completions (/v1/chat/completions):** Automatically routes requests by task phase (planning vs. execution) and logs to Redis streams.
- **Continuous Harvesting:** Logged interactions are harvested into HuggingFace-compatible JSONL datasets via `scripts/export_claude_opus_rl_dataset.py`.

6 FORMAL VERIFICATION IN LEAN 4 & BANACH FIXED-POINT AUTOPOIESIS

To guarantee that self-referential code improvement terminates without infinite refactoring loops or thermodynamic degradation, ANSE formalizes the autopoietic update loop in **Lean 4** (`formal/ANSE/BanachFixedPoint.lean`).

Theorem 1 (Autopoietic Banach Contraction). Let $\mathcal{A}$ be a complete metric space of verified AST execution graphs equipped with the metric:

$$d(A_1, A_2) = |E(A_1) - E(A_2)| + d_{\text{AST}}(A_1, A_2) \quad (7)$$

If the self-refactoring transformation $\Phi : \mathcal{A} \rightarrow \mathcal{A}$ is a strict contraction with Lipschitz constant $k < 1$:

$$d(\Phi(A_1), \Phi(A_2)) \leq k \cdot d(A_1, A_2) \quad (8)$$

then there exists a unique, minimum-energy fixed-point architecture $A^* \in \mathcal{A}$ such that $\Phi(A^*) = A^*$, and the sequence $A_{n+1} = \Phi(A_n)$ converges geometrically.

Furthermore, every process hot-swap requires the thermodynamic condition:

$$\Delta E = E(A_{\text{child}}) - E(A_{\text{parent}}) < 0 \quad (9)$$

ensuring that system entropy monotonically decreases throughout execution.

7 CONCLUSION & FUTURE OUTLOOK

The findings presented in this paper demonstrate that frontier LLMs cannot be safely evaluated through ungrounded textual benchmarks or subjective self-certification. By establishing **Physical Hardness** through zero-trust AST inspection, physical conservation invariant checking, and DPO energy margin separation, we achieve 100% verified execution across 120 PhD-level multidisciplinary problems. This paradigm shifts the foundation of autonomous AI from linguistic mimicry to verified computational physics.

REFERENCES

- [1] Y. LeCun, S. Chopra, R. Hadsell, M. Ranzato, and F. J. Huang, "A tutorial on energy-based learning," in *Predicting Structured Data*, G. Bakir et al., Eds. MIT Press, 2006.
- [2] R. Rafailov, A. Sharma, E. Mitchell, C. D. Manning, S. Ermon, and C. Finn, "Direct preference optimization: Your language model is secretly a reward model," *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 36, 2024.
- [3] M. Assran, Q. Duval, I. Caron, P. Bojanowski, P. Vincent, M. Rabbat, N. LeCun, and N. Ballas, "Self-supervised learning from images with a joint-embedding predictive architecture," in *CVPR*, 2023.
- [4] A. Bardes, J. Ponce, and Y. LeCun, "VICReg: Variance-invariance-covariance regularization for self-supervised learning," in *ICLR*, 2022.
- [5] K. Friston, "The free-energy principle: a unified brain theory?" *Nature Reviews Neuroscience*, vol. 11, no. 2, pp. 127–138, 2010.
- [6] E. Hairer, C. Lubich, and G. Wanner, *Geometric Numerical Integration: Structure-Preserving Algorithms for Ordinary Differential Equations*, 2nd ed. Springer, 2006.
- [7] Y. V. Matiyasevich, "Enumerable sets are diophantine," *Doklady Akademii Nauk SSSR*, vol. 191, no. 2, pp. 279–282, 1970.